

```
In [79]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [80]: train= pd.read_csv('Project_1_Mercedes-Benz_Greener_Manufacturing/train.csv')
test = pd.read_csv('Project_1_Mercedes-Benz_Greener_Manufacturing/test.csv')
```

```
In [81]: train.head()
```

Out[81]:

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0

5 rows × 378 columns

```
In [82]: test.head()
```

Out[82]:

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0

5 rows × 377 columns

```
In [83]: train.drop(['ID'],axis=1,inplace=True)
train.head()
```

Out[83]:

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380
0	130.81	k	v	at	a	d	u	j	o	0	...	0	0	1	0	0	0
1	88.53	k	t	av	e	d	y	l	o	0	...	1	0	0	0	0	0
2	76.26	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	0
3	80.62	az	t	n	f	d	x	l	e	0	...	0	0	0	0	0	0
4	78.02	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	0

5 rows × 377 columns

```
In [84]: test.drop(['ID'],axis=1,inplace=True)
test.head()
```

```
Out[84]:
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	X380
0	az	v	n	f	d	t	a	w	0	0	...	0	0	0	1	0	0
1	t	b	ai	a	d	b	g	y	0	0	...	0	0	1	0	0	0
2	az	v	as	f	d	a	j	j	0	0	...	0	0	0	1	0	0
3	az	l	n	f	d	z	l	n	0	0	...	0	0	0	1	0	0
4	w	s	as	c	d	y	i	m	0	0	...	1	0	0	0	0	0

5 rows × 376 columns

```
In [85]: ser = pd.Series((train[list(train.columns)[8:]].var()==0))
zero_var_cols = []
non_zero_var_cols = []
for i in range(10,500):
    if list(ser.index).count("X"+str(i))>0:
        if ser["X"+str(i)]==True:
            zero_var_cols.append("X"+str(i))
        else:
            non_zero_var_cols.append("X"+str(i))
```

/var/folders/h8/4hprg5r52wqcczhnkwfnfgt00000gn/T/ipykernel_23787/2285104892.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
ser = pd.Series((train[list(train.columns)[8:]].var()==0))
```

```
In [86]: train.drop(zero_var_cols,inplace=True,axis=1)
train.head()
```

```
Out[86]:
```

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X3
0	130.81	k	v	at	a	d	u	j	o	0	...	0	0	1	0	0	
1	88.53	k	t	av	e	d	y	l	o	0	...	1	0	0	0	0	
2	76.26	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	
3	80.62	az	t	n	f	d	x	l	e	0	...	0	0	0	0	0	
4	78.02	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	

5 rows × 365 columns

```
In [87]: test.drop(zero_var_cols,inplace=True,axis=1)
test.head()
```

```
Out[87]:
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	X379	X380
0	az	v	n	f	d	t	a	w	0	0	...	0	0	0	1	0	0
1	t	b	ai	a	d	b	g	y	0	0	...	0	0	1	0	0	0
2	az	v	as	f	d	a	j	j	0	0	...	0	0	0	1	0	0
3	az	l	n	f	d	z	l	n	0	0	...	0	0	0	1	0	0
4	w	s	as	c	d	y	i	m	0	0	...	1	0	0	0	0	0

5 rows × 364 columns

```
In [88]: sum(train.isna().any())
```

```
Out[88]: 0
```

```
In [89]: sum(test.isna().any())
```

```
Out[89]: 0
```

```
In [90]: np.unique(train.dtypes)
```

```
Out[90]: array([dtype('int64'), dtype('float64'), dtype('O')], dtype=object)
```

```
In [91]: categorical_cols = ["X0", "X1", "X2", "X3", "X4", "X5", "X6", "X8"]
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for col in categorical_cols:
    le_fit = le.fit(np.union1d(np.unique(train[col]), np.unique(test[col])))
    train[col] = le.transform(train[col])
    test[col] = le.transform(test[col])
```

```
In [92]: y_train = train[['y']]
```

```
In [93]: train.drop('y', axis=1, inplace=True)
train.head()
```

```
Out[93]:
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	X379	X380
0	37	23	20	0	3	27	9	14	0	0	...	0	0	1	0	0	0
1	37	21	22	4	3	31	11	14	0	0	...	1	0	0	0	0	0
2	24	24	38	2	3	30	9	23	0	0	...	0	0	0	0	0	0
3	24	21	38	5	3	30	11	4	0	0	...	0	0	0	0	0	0
4	24	23	38	5	3	14	3	13	0	0	...	0	0	0	0	0	0

5 rows × 364 columns

```
In [94]: test.head()
```

Out[94]:

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	X379	X380
0	24	23	38	5	3	26	0	22	0	0	...	0	0	0	1	0	0
1	46	3	9	0	3	9	6	24	0	0	...	0	0	1	0	0	0
2	24	23	19	5	3	0	9	9	0	0	...	0	0	0	1	0	0
3	24	13	38	5	3	32	11	13	0	0	...	0	0	0	1	0	0
4	49	20	19	2	3	31	8	12	0	0	...	1	0	0	0	0	0

5 rows × 364 columns

In [95]: `col_names = list(train.iloc[:, :].columns) # [0,1,2,3,4,5,6,7]`

In [96]: `from sklearn.preprocessing import normalize
from sklearn.preprocessing import StandardScaler
standardScaler = StandardScaler()
train = normalize(train)
test = normalize(test)
train = standardScaler.fit_transform(train)
test = standardScaler.transform(test)`

In [97]: `train = pd.DataFrame(train, columns=col_names)
test = pd.DataFrame(test, columns=col_names)`

In [98]: `train.head()`

Out[98]:

	X0	X1	X2	X3	X4	X5	X6	X8
0	-0.024634	0.946535	-0.249581	-1.469596	-0.590029	0.965096	0.243500	0.062037
1	-0.137802	0.674721	-0.154896	0.141177	-0.716450	1.238423	0.633395	-0.003078
2	-1.218495	0.802914	0.930978	-0.718179	-0.921356	0.953092	0.009931	0.835999
3	-1.085984	0.689465	1.159887	0.559218	-0.693142	1.165315	0.653477	-1.124608
4	-0.921743	1.071072	1.443609	0.745475	-0.410282	-0.240028	-1.210723	0.030858

5 rows × 364 columns

In [99]: `test.head()`


Out[99]:

	X0	X1	X2	X3	X4	X5	X6	X8
0	-1.134464	0.817244	1.076139	0.504240	-0.776636	0.716000	-2.001191	0.848176
1	0.893821	-0.953582	-1.126072	-1.469596	-0.387743	-0.755110	-0.409663	1.412119
2	-0.288188	1.827053	0.303740	1.463953	0.680840	-1.699863	1.139399	-0.102898
3	-1.086233	-0.053212	1.159457	0.558936	-0.693571	1.355901	0.653108	-0.104632
4	0.422747	0.416328	-0.535453	-0.736904	-0.992447	0.973211	-0.258074	-0.349918

5 rows × 364 columns

```
In [372... from sklearn.decomposition import PCA
pca = PCA()
x_train = pca.fit_transform(train)
x_test = pca.transform(test)
plt.figure(figsize=(15,5))
plt.style.use("ggplot")
plt.plot(pca.explained_variance_, marker='o')
plt.xlabel("Eigenvalue number")
plt.ylabel("Eigenvalue size")
plt.title("Scree Plot")
```

Out[372]: Text(0.5, 1.0, 'Scree Plot')



From above CDF graph we can see that maximum variance can be explained using 25 principal components (the elbow point)

```
In [425... pca = PCA(n_components=25)
x_train = pca.fit_transform(train)
x_test = pca.transform(test)
```

```
In [426... import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import accuracy_score
from sklearn.metrics import r2_score
```

```
In [427... X_train, X_test, Y_train, Y_test = train_test_split(x_train, y_train, train_s
```

```
In [428... regressor = xgb.XGBRegressor(n_estimators=20, reg_lambda=40, gamma=600, max_dep
regressor.fit(X_train, Y_train)
pd.DataFrame(regressor.feature_importances_.reshape(1, -1))
```

```
Out[428]:
```

	0	1	2	3	4	5	6	7	
0	0.024092	0.226885	0.150361	0.02353	0.044362	0.049162	0.022303	0.019392	0.02020


1 rows x 25 columns

```
In [429... Y_pred = regressor.predict(X_test)
print("MSE: " + str(mean_squared_error(Y_test, Y_pred)) + " R2-Score: " + str(r
```

MSE: 87.12105784499767 R2-Score: 0.4800502881611578

```
In [430... plt.figure(figsize=(20,10))
plt.scatter(range(0,len(Y_pred)),Y_pred,c='black',s=10)
plt.scatter(range(0,len(Y_test)),Y_test,c='red',s=10)
```

Out[430]: <matplotlib.collections.PathCollection at 0x7fb828c86dc0>




Testing the actual test data provided

```
In [432... y_test = regressor.predict(x_test)
```

```
In [436... plt.scatter(range(0,len(y_test)),y_test,c='red',s=10)
```

Out[436]: <matplotlib.collections.PathCollection at 0x7fb82805d8e0>



In []: